

twitter-sentiment-analysis

January 5, 2024

```
[49]: #Importing Necessary Libraries
import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score
from imblearn.over_sampling import SMOTE

from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier

from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from wordcloud import WordCloud, STOPWORDS
import re
import warnings
warnings.filterwarnings('ignore')
```

1 Data Loading

```
[50]: train_df = pd.read_csv('/kaggle/input/twittersentimentdata/train.csv')
test_df = pd.read_csv('/kaggle/input/twittersentimentdata/test.csv')
```

```
[51]: train_df.shape
```

```
[51]: (31962, 3)
```

```
[52]: train_df.duplicated().sum()
```

```
[52]: 0
```

```
[53]: train_df.dtypes
```

```
[53]: id          int64  
      label      int64  
      tweet      object  
      dtype: object
```

```
[54]: train_df.isnull().sum()
```

```
[54]: id          0  
      label      0  
      tweet      0  
      dtype: int64
```


```
[55]: test_df.isnull().sum()
```

```
[55]: id          0  
      tweet      0  
      dtype: int64
```


```
[56]: # Plotting Word Clouds  
      stopwords = set(STOPWORDS)  
      stopwords.add('user')  
  
      def plot_wordcloud(tweets, title):  
          wordcloud = WordCloud(width=800, height=800, background_color='white',  
                                ↪stopwords=stopwords, min_font_size=10).generate(tweets)  
          plt.figure(figsize=(14, 6), facecolor=None)  
          plt.imshow(wordcloud)  
          plt.axis("off")  
          plt.title(title, fontdict={'fontsize': 20})  
          plt.show()
```

```
[57]: negative_tweets = train_df['tweet'][train_df['label'] == 1].to_string()  
      positive_tweets = train_df['tweet'][train_df['label'] == 0].to_string()  
  
      plot_wordcloud(negative_tweets, 'Negative Tweets')  
      plot_wordcloud(positive_tweets, 'Positive Tweets')
```

Negative Tweets



Positive Tweets



2 Feature Engineering

```
[58]: # Feature Engineering
train_df_fe = train_df.copy()
train_df_fe['tweet_length'] = train_df_fe['tweet'].str.len()
train_df_fe['num_hashtags'] = train_df_fe['tweet'].str.count('#')
train_df_fe['num_exclamation_marks'] = train_df_fe['tweet'].str.count('!')
train_df_fe['num_question_marks'] = train_df_fe['tweet'].str.count('\?')
train_df_fe['total_tags'] = train_df_fe['tweet'].str.count('@')
train_df_fe['num_punctuations'] = train_df_fe['tweet'].str.count('[.,,:;]')
train_df_fe['num_words'] = train_df_fe['tweet'].apply(lambda x: len(x.split()))
train_df_fe.head()
```

```
[58]:    id  label          tweet  tweet_length \
0    1      0  @user when a father is dysfunctional and is s...      102
1    2      0  @user @user thanks for #lyft credit i can't us...      122
```

2	3	0		bihday your majesty	21
3	4	0	#model	i love u take with u all the time in ...	86
4	5	0		factsguide: society now #motivation	39

	num_hashtags	num_exclamation_marks	num_question_marks	total_tags	\
0	1	0	0	1	
1	3	0	0	2	
2	0	0	0	0	
3	1	3	0	0	
4	1	0	0	0	


	num_punctuations	num_words
0	1	18
1	1	19
2	0	3
3	0	14
4	1	4

```
[59]: # Visualizing Relationship of Engineered Features with Sentiments
features = ['tweet_length', 'num_hashtags', 'num_exclamation_marks',
           ↪ 'num_question_marks', 'total_tags', 'num_punctuations', 'num_words']

# Check if train_df_fe has the expected columns
if set(features).issubset(train_df_fe.columns):
    plt.figure(figsize=(12, 16))
    colors = ['skyblue', 'salmon']

    for i, feature in enumerate(features, 1):
        plt.subplot(4, 2, i)
        sns.histplot(train_df_fe[train_df_fe.label == 0][feature],
           ↪ label='Positive', kde=True, color=colors[0])
        sns.histplot(train_df_fe[train_df_fe.label == 1][feature],
           ↪ label='Negative', kde=True, color=colors[1])
        plt.legend()

    plt.tight_layout()
    plt.show()
```



3 Data Preprocessing

```
[60]: # Data Preprocessing
X = train_df.drop(columns=['label'])
y = train_df['label']
test = test_df
```

```
[61]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

3.1 Text Normalization

```
[62]: def tokenize_and_clean(text):
    lowered = text.lower()
    cleaned = re.sub('@user', '', lowered)
    tokens = word_tokenize(cleaned)
    filtered_tokens = [token for token in tokens if re.match(r'\w{1,}', token)]
    stemmer = PorterStemmer()
    stems = [stemmer.stem(token) for token in filtered_tokens]
    return stems
```

3.2 Vectorization

```
[63]: # TF-IDF Vectorization
tfidf_vectorizer = TfidfVectorizer(tokenizer=tokenize_and_clean,
↳ stop_words='english')
X_train_tweets_tfidf = tfidf_vectorizer.fit_transform(X_train['tweet'])
X_test_tweets_tfidf = tfidf_vectorizer.transform(X_test['tweet'])
X_tweets_tfidf = tfidf_vectorizer.fit_transform(X['tweet'])
test_tweets_tfidf = tfidf_vectorizer.transform(test['tweet'])
```

3.3 SMOTE

```
[64]: # Class Imbalance Check Before SMOTE
plt.figure(figsize=(12, 6))

# Colors for the pie charts
colors_before_smote = ['#66b3ff', '#99ff99']
colors_after_smote = ['#ff9999', '#66b3ff']

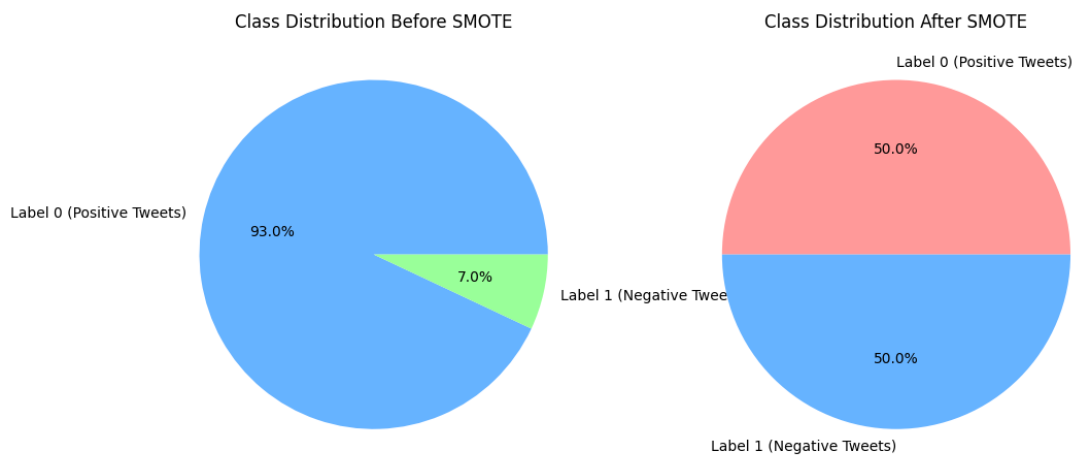
labels = ['Label 0 (Positive Tweets)', 'Label 1 (Negative Tweets)']

# Plotting before SMOTE
plt.subplot(1, 2, 1)
plt.pie(y_train.value_counts(), labels=labels, autopct='%0.1f%%',
↳ colors=colors_before_smote)
```

```
plt.title('Class Distribution Before SMOTE')

# SMOTE to deal with the class imbalance
smote = SMOTE()
X_train_smote, y_train_smote = smote.fit_resample(X_train_tweets_tfidf, y_train.
↪values)

# Plotting after SMOTE
plt.subplot(1, 2, 2)
plt.pie(pd.value_counts(y_train_smote), labels=labels, autopct='%0.1f%%',
↪colors=colors_after_smote)
plt.title('Class Distribution After SMOTE')
plt.show()
```



4 ML Model

```
[65]: # Functions to print scores
def training_scores(y_act, y_pred):
    acc = round(accuracy_score(y_act, y_pred), 3)
    f1 = round(f1_score(y_act, y_pred), 3)
    print(f'Training Scores: Accuracy={acc}, F1-Score={f1}')

def validation_scores(y_act, y_pred):
    acc = round(accuracy_score(y_act, y_pred), 3)
    f1 = round(f1_score(y_act, y_pred), 3)
    print(f'Validation Scores: Accuracy={acc}, F1-Score={f1}')

[66]: # Machine Learning Modeling
def train_and_evaluate(model, X_train, y_train, X_test, y_test):
    model.fit(X_train, y_train)
```



```

y_train_pred = model.predict(X_train)
y_test_pred = model.predict(X_test)

training_scores(y_train, y_train_pred)
validation_scores(y_test, y_test_pred)

```

```

[67]: # Logistic Regression
lr = LogisticRegression()
train_and_evaluate(lr, X_train_smote, y_train_smote, X_test_tweets_tfidf,
    ↪y_test)

```

Training Scores: Accuracy=0.974, F1-Score=0.975
 Validation Scores: Accuracy=0.924, F1-Score=0.601

```

[68]: # Naive Bayes Classifier
mnb = MultinomialNB()
train_and_evaluate(mnb, X_train_smote, y_train_smote, X_test_tweets_tfidf,
    ↪y_test)

```

Training Scores: Accuracy=0.966, F1-Score=0.967
 Validation Scores: Accuracy=0.921, F1-Score=0.609

```

[69]: # Random Forest Classifier
rf = RandomForestClassifier()
train_and_evaluate(rf, X_train_smote, y_train_smote, X_test_tweets_tfidf,
    ↪y_test)

```

Training Scores: Accuracy=1.0, F1-Score=1.0
 Validation Scores: Accuracy=0.955, F1-Score=0.648

```

[70]: # Extreme Gradient Boosting Classifier
xgb = XGBClassifier(objective='binary:logistic', eval_metric='logloss')
train_and_evaluate(xgb, X_train_smote, y_train_smote, X_test_tweets_tfidf,
    ↪y_test)

```

Training Scores: Accuracy=0.941, F1-Score=0.938
 Validation Scores: Accuracy=0.942, F1-Score=0.597

5 Hyperparameter Tuning

```

[71]: rf_tuned = RandomForestClassifier(criterion='entropy',
    max_samples=0.8,
    min_samples_split=10,
    random_state=0)
train_and_evaluate(rf_tuned, X_train_smote, y_train_smote, X_test_tweets_tfidf,
    ↪y_test)

```

Training Scores: Accuracy=0.999, F1-Score=0.999
Validation Scores: Accuracy=0.957, F1-Score=0.682

```
[72]: xgb_tuned = XGBClassifier(objective='binary:logistic',  
                               eval_metric='logloss',  
                               learning_rate=0.8,  
                               max_depth=20,  
                               gamma=0.6,  
                               reg_lambda=0.1,  
                               reg_alpha=0.1)  
train_and_evaluate(xgb_tuned, X_train_smote, y_train_smote,  
                  ↪X_test_tweets_tfidf, y_test)
```

Training Scores: Accuracy=0.998, F1-Score=0.998
Validation Scores: Accuracy=0.954, F1-Score=0.65

Your support with an **upvote** would be greatly valued.